

CONTACT INFORMATION	New York University Center for Data Science 60 5th Avenue, 7th Floor, Room 602 New York, NY 10011 USA	Phone +1-917-696-4069 <i>E-mail:</i> djellel@nyu.edu <i>WWW:</i> https://dedcode.github.io
RESEARCH INTERESTS	Crowdsourcing Systems, Crowd Dynamics, Data Management Systems, Distributed Systems, Data Mining, Machine Learning, Knowledge Extraction	
EDUCATION	University of Fribourg , Fribourg, Switzerland 2011 to 2015 Ph.D., Computer Science (Best CS Thesis Award) – June 2015 <ul style="list-style-type: none">• Dissertation title: “Quality of Service in Crowd-Powered Systems”• Adviser: Professor Philippe Cudré-Mauroux University of Louisiana , Lafayette, USA 2009 to 2011 M.S., Computer Science (with Honors) – May 2011 <ul style="list-style-type: none">• Adviser: Professor Vijay V. Raghavan• Topic: Action Rules Mining• Project: “Frequent Association Action Rules Mining Using FP-Trees” University of Sciences and Technology H. Boumediene , Algiers, Algeria 1999 to 2004 Dipl.-Ing., Informatics – December 2004 <ul style="list-style-type: none">• Track: Information Systems• Project: Modeling a Medical Center Information System with UML	
PROFESSIONAL EMPLOYMENT	Moore-Sloan Data Science Fellow , New York University June 2017 to present NYU Center for Data Science <ul style="list-style-type: none">• Research in crowd population estimation and dynamics; crowd behavior analysis in Wikidata community; Wikidata-based data integration; Wikidata graph evolution prediction; Crowdsourcing sport video annotation Postdoctoral Researcher , University of Fribourg (Switzerland) September 2015 to May 2017 eXascale Infolab <ul style="list-style-type: none">• Research in query expansion (tailored to job search), entity linking and co-reference resolution• Supervision of PhD and Master students• Supervisor: Professor Philippe Cudré-Mauroux Research Assistant , University of Fribourg (Switzerland) May 2011 to June 2015 eXascale Infolab <ul style="list-style-type: none">• Research in crowd-powered systems: Human intelligence task scheduling, task routing, and worker retention.• Research in crowdsourcing-based entity linking and instance matching• Supervision of Bachelor and Master students• Supervisor: Professor Philippe Cudré-Mauroux Intern , Microsoft Research CISL, Mountain View, CA May to August 2013 Cloud and Information Services Lab <ul style="list-style-type: none">• Project: “Reservation based scheduling with Hadoop YARN”• Supervisor: Dr. Carlo Curino Student Developer , Google Summer of Code Program May to August 2010 Drizzle DBMS <ul style="list-style-type: none">• Project: “A query cache plugin for Drizzle DBMS based on Memcached”• Supervisor: Toru Maesaka	

	<i>Travel Grant – HCOMP Doctoral Consortium</i> <ul style="list-style-type: none"> • Project: “Scalable Human-based Computation” • Grant value: 700 USD <i>University of Louisiana at Lafayette Honors, for maintaining a GPA of 4.0</i> <i>Fulbright Exchange Student Grant</i>	October 2014 2011 2009-2011
INVITED TALKS	<ul style="list-style-type: none"> • Djellel Difallah. “The Dynamics of Crowdsourcing Platforms: The Case of Amazon Mechanical Turk and Wikidata” (Seminar) NYU Center for Data Science. NY, USA • Ujwal Gadiraju, Gianluca Demartini, Djellel Difallah and Michele Catasta. “Using Crowdsourcing Effectively for Social Media Research” (Tutorial) ICWSM’16 Cologne, Germany • Ujwal Gadiraju, Gianluca Demartini, Djellel Difallah and Michele Catasta. “Using Crowdsourcing Effectively for Social Media Research” (Tutorial) WebSci’16 Hannover, Germany • Djellel Difallah. “Quality and Performance Optimizations in Microtask Crowdsourcing” MIT CSAIL, Boston, MA • Djellel Difallah. “Quality and Performance Optimizations in Microtask Crowdsourcing” Parallel Data Lab, CMU, Pittsburgh, PA • Djellel Difallah. “Extending The OLTP-Bench Framework for Big Data Systems” (Workshop) Workshop on Big Data Benchmarking, Postdam, Germany • Djellel Difallah. “Resource Management and Planning for Hadoop” Microsoft Research. Mountain View, CA • Djellel Difallah. “Mechanical Cheat: Spamming Schemes and Adversarial Techniques on Crowdsourcing Platforms” (Workshop) CrowdSearch 2012 workshop at WWW, Lyon, France 	September 2017 May 2016 May 2016 November 2014 October 2014 August 2014 August 2013 April 2012
PROFESSIONAL SERVICE	Program Committee Member <ul style="list-style-type: none"> • 24th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD 2018) • 12th International AAAI Conference On WEB and Social Media (ICWSM 2018) • 27th International World Wide Web Conference (WWW 2018) • 23rd ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD 2017) • 26th ACM Conference on Information and Knowledge Management (CIKM 2017) • 26th International World Wide Web Conference (WWW 2017) • 16th International Conference on Web Engineering (ICWE 2016) External Reviewer <ul style="list-style-type: none"> • IEEE International Conference on Data Engineering, ICDE (2016) • ACM Special Interest Group on Management of Data, SIGMOD (2016) • ACM Conference on Web Science, WebSci (2016) • ACM Symposium on Cloud Computing, SoCC (2016) • Very Large Database, VLDB (2014, 2016) • International Conference on Extending Database Technology, EDBT (2013) • Association for the Advancement of Artificial Intelligence, AAAI (2014, 2015) • IEEE BigData (2013) 	
OPEN SOURCE SERVICE	<i>OLTP-Bench</i> (Lead Contributor 2012-2015) <ul style="list-style-type: none"> • Database Benchmarking Framework • https://github.com/oltpbenchmark <i>OpenTurk</i> (Lead Contributor 2013-2014) <ul style="list-style-type: none"> • Chrome extension for Mturk HIT notification • https://github.com/openturk 	

Apache Hadoop YARN (Contributor 2013)

- Research and development on the scheduling system for Hadoop 2.0
- Code Repository: <https://issues.apache.org/jira/browse/YARN-1051>
- Various bug fixes in the scheduler

SciDB (Contributor 2011-2013)

- Array Database Management System
- Implementation of a science benchmark (SSDB)

Drizzle (Contributor 2010)

- Database Management System with a micro-kernel architecture
- Code Repository: <https://code.launchpad.net/~dedzone/drizzle/query-cache-soc>

Apache Mahout (Contributor 2010)

- Scalable machine-learning algorithms
- Bug fixes in K-Means implementation

ADVISING AND
MENTORING

University Students

- **Alexander Striffeler**, Master Student, Computer Science 2015–
Thesis topic: “Benchmarking Big Data Infrastructures”
Co-advised by: Philippe Cudré-Mauroux and Mourad Khayati
- **Marwa Bouzeyane**, Master Student, Computer Science 2015–
Thesis topic: “Grammatical Errors Detection Based On N-gram Analysis”
Co-advised by: Philippe Cudré-Mauroux and Mourad Khayati
- **Stefan Nüesch**, Master Student, Computer Science 2014
Thesis topic: “Real-Time Anomaly Detection in Water Distribution Networks using Spark Streaming”
Co-advised by: Philippe Cudré-Mauroux
- **Phokham Nonava**, Master Student, Computer Science 2014
Thesis topic: “HDFS Blocks Placement Strategy”
Co-advised by: Philippe Cudré-Mauroux, Benoit Perroud and Martin Grund
- **Simpal Kumar**, Master Student, Computer Science 2014
Thesis topic: “Real Time Data Analysis for Water Distribution Network using Storm”
Co-advised by: Philippe Cudré-Mauroux
- **Dani Rotzetter**, Master Student, Computer Science 2014
Thesis topic: “Crowd-Flow Designer: An Open-Source Toolkit to Design and Run Complex Crowd-Sourced Tasks”
Co-advised by: Philippe Cudré-Mauroux and Gianluca Demartini
- **Victor Felder**, Bachelor Student, Computer Science 2013
Thesis topic: “Openturk: An Implementation of the Pick-A-Crowd Architecture”

High School Students

- **Job-Info Day Fribourg**. Introduction to the computer science career track for high school students (2013-2017).
- **Cyber Camp**. Instructor at the computer science summer camp for high school students (2012-2014).

CONFERENCE
PUBLICATIONS

- [1] J. Plu, R. Prokofyev, T. Tonon, P. Cudré-Mauroux, **D. Difallah**, R. Troncy, and G. Rizzo. Sanaphor++: Combining deep neural networks with semantics for coreference resolution. LREC, 2018
- [2] **D. Difallah**, E. Filatova, and P. Ipeirotis. Demographics and dynamics of mechanical turk workers. In *Proceedings of the Eleventh ACM International Conference on Web Search and Data Mining*, WSDM ’18, pages 135–143, New York, NY, USA, 2018. ACM
- [3] R. Prokofyev, M. Luggen, **D. Difallah**, and P. Cudré-Mauroux. Swisslink: High-precision,

context-free entity linking exploiting unambiguous labels. In *Proceedings of the 13th International Conference on Semantic Systems*. ACM, 2017

- [4] A. Tonon, V. Felder, **D. Difallah**, and P. Cudré-Mauroux. Voldemortkg: Mapping schema.org and web entities to linked open data. In *The Semantic Web - ISWC 2016 - 15th International Semantic Web Conference, Kobe, Japan, October 17-21, 2016, Proceedings, Part II*, pages 220–228, 2016
- [5] R. Prokofyev, A. Tonon, M. Luggen, L. Vouilloz, **D. Difallah**, and P. Cudré-Mauroux. SANAPHOR: ontology-based coreference resolution. In *The Semantic Web - ISWC 2015 - 14th International Semantic Web Conference, Bethlehem, PA, USA, October 11-15, 2015, Proceedings, Part I*, pages 458–473, 2015
- [6] **D. Difallah**, M. Catasta, G. Demartini, P. G. Ipeirotis, and P. Cudré-Mauroux. The dynamics of micro-task crowdsourcing: The case of amazon mturk. In *Proceedings of the 24th International Conference on World Wide Web, WWW 2015, Florence, Italy, May 18-22, 2015*, pages 238–247, 2015
- [7] M. Catasta, A. Tonon, **D. Difallah**, G. Demartini, K. Aberer, and P. Cudré-Mauroux. Hippocampus: answering memory queries using transactive search. In *23rd International World Wide Web Conference, WWW '14, Seoul, Republic of Korea, April 7-11, 2014, Companion Volume*, pages 535–540, 2014
- [8] M. Catasta, A. Tonon, **D. Difallah**, G. Demartini, K. Aberer, and P. Cudré-Mauroux. Transactivedb: Tapping into collective human memories. *PVLDB*, 7(14):1977–1980, 2014
- [9] C. Curino, **D. Difallah**, C. Douglas, S. Krishnan, R. Ramakrishnan, and S. Rao. Reservation-based scheduling: If you're late don't blame us! In *Proceedings of the ACM Symposium on Cloud Computing, Seattle, WA, USA, November 03 - 05, 2014*, pages 2:1–2:14, 2014
- [10] **D. Difallah**, M. Catasta, G. Demartini, and P. Cudré-Mauroux. Scaling-up the crowd: Micro-task pricing schemes for worker retention and latency improvement. In *Proceedings of the Seconf AAAI Conference on Human Computation and Crowdsourcing, HCOMP 2014, November 2-4, 2014, Pittsburgh, Pennsylvania, USA*, 2014
- [11] G. Demartini, **D. Difallah**, and P. Cudré-Mauroux. Large-scale linked data integration using probabilistic reasoning and crowdsourcing. *VLDB J.*, 22(5):665–687, 2013
- [12] **D. Difallah**, A. Pavlo, C. Curino, and P. Cudré-Mauroux. Oltp-bench: An extensible testbed for benchmarking relational databases. *PVLDB*, 7(4):277–288, 2013
- [13] **D. Difallah**, G. Demartini, and P. Cudré-Mauroux. Pick-a-crowd: tell me what you like, and i'll tell you what to do. In *22nd International World Wide Web Conference, WWW '13, Rio de Janeiro, Brazil, May 13-17, 2013*, pages 367–374, 2013
- [14] **D. Difallah**, P. Cudré-Mauroux, and S. A. McKenna. Scalable anomaly detection for smart city infrastructure networks. *IEEE Internet Computing*, 17(6):39–47, 2013
- [15] G. Demartini, **D. Difallah**, and P. Cudré-Mauroux. Zencrowd: leveraging probabilistic reasoning and crowdsourcing techniques for large-scale entity linking. In *Proceedings of the 21st World Wide Web Conference 2012, WWW 2012, Lyon, France, April 16-20, 2012*, pages 469–478, 2012
- [16] G. Demartini, **D. Difallah**, and P. Cudré-Mauroux. Large-scale linked data integration using probabilistic reasoning and crowdsourcing. *VLDB J.*, 22(5):665–687, 2013
- [17] **D. Difallah**, P. Cudré-Mauroux, and S. A. McKenna. Scalable anomaly detection for smart city infrastructure networks. *IEEE Internet Computing*, 17(6):39–47, 2013

- BOOKS [18] G. Demartini, **D. Difallah**, U. Gadiraju, M. Catasta, et al. An introduction to hybrid human-machine information systems. *Foundations and Trends® in Web Science*, 7(1):1–87, 2017
- BOOK CHAPTERS [19] **D. Difallah**, P. Cudré-Mauroux, S. McKenna, and D. Fasel. Skalierbar anomalien erkennen für smart city infrastrukturen. In *Big Data*, pages 289–299. Springer, 2016
- WORKSHOPS AND DEMONSTRATIONS [20] D. V. Aken, **D. Difallah**, A. Pavlo, C. Curino, and P. Cudré-Mauroux. Benchpress: Dynamic workload control in the oltp-bench testbed. In *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data, Melbourne, Victoria, Australia, May 31 - June 4, 2015*, pages 1069–1073, 2015
- [21] C. Curino, **D. Difallah**, A. Pavlo, and P. Cudré-Mauroux. Benchmarking oltp/web databases in the cloud: the oltp-bench framework. In *Proceedings of the Fourth International Workshop on Cloud Data Management, CloudDB 2012, Maui, HI, USA, October 29, 2012*, pages 17–20, 2012
- [22] **D. Difallah**, G. Demartini, and P. Cudré-Mauroux. Mechanical cheat: Spamming schemes and adversarial techniques on crowdsourcing platforms. In *Proceedings of the First International Workshop on Crowdsourcing Web Search, Lyon, France, April 17, 2012*, pages 26–30, 2012
- [23] P. Cudré-Mauroux, G. Demartini, **D. Difallah**, A. Mostafa, V. Russo, and M. Thomas. A Demonstration of DNS3: a Semantic-Aware DNS Service. Citeseer, 2011
- [24] **D. Difallah**, R. G. Benton, V. V. Raghavan, and T. Johnsten. FAARM: frequent association action rules mining using fp-tree. In *Data Mining Workshops (ICDMW), 2011 IEEE 11th International Conference on, Vancouver, BC, Canada, December 11, 2011*, pages 398–404, 2011
- [25] **D. Difallah**, P. Cudré-Mauroux, I. Stratos, and Y. Zhang. SSDB Benchmark: Implementations and Results. XLDB, 2011