

Djellel Difallah

Address: 888 Main Street Apt. 112

New York, NY 10044 – USA

Email: djellel@nyu.edu, Phone: +1 (917) 696-4069

Webpage: <https://dedcode.github.io/>

Research Interests

My main research interests revolve around human computation and its application to data management, data processing, machine learning, natural language processing and knowledge extraction. Specifically, I investigate novel human machine learning techniques, Bayesian models for truth inference, task scheduling optimization problems and analysis of crowdsourcing marketplaces.

Education

- | | |
|-----------|---|
| 2011–2015 | PhD in Computer Science, University of Fribourg – Switzerland. <ul style="list-style-type: none">– Dissertation title: “Quality of Service in Crowd-Powered Systems”.– Received the best thesis award in Computer Science (University of Fribourg 2015). |
| 2009–2011 | MSc in Computer Science, University of Louisiana at Lafayette – USA. <ul style="list-style-type: none">– Fulbright Scholar.– Received honors for maintaining a GPA of 4.0 four semesters. |
| 1999–2004 | Dipl.-Ing. in Informatics, University of Science and Technology of Algiers – Algeria. |

Employment history

- | | |
|--------------|--|
| Since 2017 | Moore-Sloan Data Science Fellow (New York University). <ul style="list-style-type: none">– Current research focuses on studying crowdsourcing platforms, their dynamics and evolution both from a worker and economic perspective and based on public activity logs.– Pursue a research plan on combining machine learning and crowdsourcing in the context of data science. |
| 2015–05/2017 | Senior Research Scientist at eXascale Infolab (University of Fribourg). <ul style="list-style-type: none">– Postdoctoral research on active learning with the crowd, and truth inference from crowdsourced data; applied to data analysis and information retrieval projects.– Contribute to research projects on entity linking, entity resolution, and query expansion.– Lecturer for the course “Big Data Infrastructures” and the seminar “Data Science”. |
| 2011–2015 | Research Assistant at eXascale Infolab (University of Fribourg). <ul style="list-style-type: none">– Contributed to various projects in the group e.g., anomaly detection (stream processing), array data management (SciDB), human memory-based information systems (TransactiveDB), entity linking and coreference resolution (semantic Web and NLP).– Teaching assistant for the social computing class.– Supervise master students working on stream data processing solutions for anomaly detection in Water Distribution Networks. |
| 2013 | Internship at Cloud and Information Services Lab (Microsoft Research). <ul style="list-style-type: none">– Project: “Reservation based scheduling with Hadoop YARN”.– Repo: https://issues.apache.org/jira/browse/YARN-1051 |
| 2010 | Student Developer at Google Summer of Code Program. <ul style="list-style-type: none">– Project: “A query cache plugin for Drizzle DBMS based on Memcached”. |

	– Repo: https://code.launchpad.net/~dedzone/drizzle/query-cache-soc
2006–2009	Information Management Engineer at Schlumberger. – Database administration (Oracle). – Dashboard and KPI development. – SQL reports and query optimization.
2005–2006	Engineer at EEPAD Internet Services Provider. – System administration of the authentication system platform (RADIUS) for both line and wireless clients. – Developed an automatic provisioning solution to synchronize the information system and the deployed ADSL infrastructure using SNMP.

Teaching Experience

	(<i>Swiss Joint Master of Science in Computer Science</i> http://mcs.unibn.ch/)
Lecturer	Big Data Infrastructures (Graduate – Fall 2016). – Main Topics: NOSQL systems (Dynamo, Cassandra, MongoDB), Hadoop Map/Reduce, YARN, Spark and Kafka. – Co-Instructor: Philippe Cudré-Mauroux.
Lecturer	Data Science Seminar (Graduate – Fall 2015). – Topics: Selected papers focusing on time series analysis and crowdsourcing. – Co-Instructor: Mourad Khayati.
TA	Social Computing (Graduate – Fall semester from 2011 to 2014). – Instructors: Philippe Cudré-Mauroux and Gianluca Demartini.

Professional Activities

PC member	– WWW(2017), CIKM(2017), KDD(2017), CHI(2017), ICWE(2016), ESWC(2015,2016).
Reviewer	– I routinely review journal publications for TKDE, TSC, VLDBJ, SWJ as well as conference papers for WWW, VLDB, SIGMOD, ICDE, SoCC, ISWC.

Grants and Other Funding

2016	EU H2020 ICT-14. – Project track: “Big Data PPP: Cross-sectorial and Cross-lingual Data Integration and Experimentation”. – Project title: “FashionBrain: Understanding Europe’s Fashion Data Universe”. – Three academic institutions and three industry partners. – Grant value: 2.9M EUR total cost; University of Fribourg share: 690’000 EUR.
2016	Swiss Commission of Technology and Innovation. – Project title: “Query Expansion: Deep Learning and Crowdsourcing”. – A technology transfer project with one industry partner (job aggregator). – Grant value: 210’000 CHF.
2015	European Science Foundation (Research Visit Grant). – Host: Dr. Gianluca Demartini University of Sheffield. – Theme: “Measuring Human Behaviour in Search Evaluation Micro Tasks”. – Grant value: 1’350 EUR.
2014	HCOMP Doctoral Consortium (Travel Grant). – Theme: “Scalable Human-based Computation”. – Travel grant value: 700 USD.

Talks and Tutorials

- Introduction to Big Data* (CAS Business Intelligence and Analytics).
 - At Lucerne University of Applied Sciences. Lucerne, Switzerland – August 2016.
- Using Crowdsourcing Effectively for Social Media Research* (tutorial).
 - At ICWSM. Cologne, Germany – May 2016.
 - At WebSci. Hannover, Germany – May 2016.
 - Together With Ujwal Gadiraju, Gianluca Demartini, and Michele Catasta.
- Quality and Performance Optimizations in Microtask Crowdsourcing*.
 - At MIT CSAIL. Boston, MA – November 2014.
 - At CMU. Pittsburgh, PA – October 2014.
- Extending The OLTP-Bench Framework for Big Data Systems*.
 - At the Workshop on Big Data Benchmarking. Postdam, Germany – August 2014.
- Resource Management and Planning for Hadoop*.
 - At Microsoft Cloud and Information Services Lab. Mountain View, CA – August 2013.

Opensource Projects

- | | |
|-------------|--|
| Lead | <ul style="list-style-type: none">– OLTP-Bench: https://github.com/oltpbenchmark– Openturk: https://github.com/openturk |
| Contributor | <ul style="list-style-type: none">– Apache Hadoop YARN (Bug fixes and reservation based scheduling PoC).– SciDB (I implemented the <i>SSDB</i> astronomy benchmark).– Drizzle (I implemented the query cache plugin).– Apache Mahout (Bug fixes). |

Spoken Languages

- Fluent in French, English, and Arabic.

List of Publications

1. Prokofyev, R., Luggen, M., **Difallah, D. E.** & Cudré-Mauroux, P. *SwissLink: High-Precision, Context-Free Entity Linking Exploiting Unambiguous Labels* in *Proceedings of the 13th International Conference on Semantic Systems* (2017).
2. **Difallah, D. E.**, Demartini, G. & Cudré-Mauroux, P. *Scheduling Human Intelligence Tasks in Multi-Tenant Crowd-Powered Systems* in *Proceedings of the 25th International Conference on World Wide Web, WWW 2016, Montreal, Canada, April 11 - 15, 2016* (2016), 855–865.
3. Tonon, A., Felder, V., **Difallah, D. E.** & Cudré-Mauroux, P. *VoldemortKG: Mapping schema.org and Web Entities to Linked Open Data* in *The Semantic Web - ISWC 2016 - 15th International Semantic Web Conference, Kobe, Japan, October 17-21, 2016, Proceedings, Part II* (2016), 220–228.
4. Aken, D. V., **Difallah, D. E.**, Pavlo, A., Curino, C. & Cudré-Mauroux, P. *BenchPress: Dynamic Workload Control in the OLTP-Bench Testbed* in *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data, Melbourne, Victoria, Australia, May 31 - June 4, 2015* (2015), 1069–1073.
5. Prokofyev, R., Tonon, A., Luggen, M., Vouilloz, L., **Difallah, D. E.** & Cudré-Mauroux, P. *SANAPHOR: Ontology-Based Coreference Resolution* in *The Semantic Web - ISWC 2015 - 14th International Semantic Web Conference, Bethlehem, PA, USA, October 11-15, 2015, Proceedings, Part I* (2015), 458–473.
6. **Difallah, D. E.**, Catasta, M., Demartini, G., Ipeirotis, P. G. & Cudré-Mauroux, P. *The Dynamics of Micro-Task Crowdsourcing: The Case of Amazon MTurk* in *Proceedings of the 24th International Conference on World Wide Web, WWW 2015, Florence, Italy, May 18-22, 2015* (2015), 238–247.

7. Catasta, M., Tonon, A., **Difallah, D. E.**, Demartini, G., Aberer, K. & Cudré-Mauroux, P. *Hippocampus: answering memory queries using transactive search* in *23rd International World Wide Web Conference, WWW '14, Seoul, Republic of Korea, April 7-11, 2014, Companion Volume* (2014), 535–540.
8. Catasta, M., Tonon, A., **Difallah, D. E.**, Demartini, G., Aberer, K. & Cudré-Mauroux, P. TransactiveDB: Tapping into Collective Human Memories. *PVLDB* **7**, 1977–1980 (2014).
9. Curino, C., **Difallah, D. E.**, Douglas, C., Krishnan, S., Ramakrishnan, R. & Rao, S. *Reservation-based Scheduling: If You're Late Don't Blame Us!* in *Proceedings of the ACM Symposium on Cloud Computing, Seattle, WA, USA, November 03 - 05, 2014* (2014), 2:1–2:14.
10. **Difallah, D. E.**, Catasta, M., Demartini, G. & Cudré-Mauroux, P. *Scaling-Up the Crowd: Micro-Task Pricing Schemes for Worker Retention and Latency Improvement* in *Proceedings of the Second AAAI Conference on Human Computation and Crowdsourcing, HCOMP 2014, November 2-4, 2014, Pittsburgh, Pennsylvania, USA* (2014).
11. Demartini, G., **Difallah, D. E.** & Cudré-Mauroux, P. Large-scale linked data integration using probabilistic reasoning and crowdsourcing. *VLDB J.* **22**, 665–687 (2013).
12. **Difallah, D. E.**, Pavlo, A., Curino, C. & Cudré-Mauroux, P. OLTP-Bench: An Extensible Testbed for Benchmarking Relational Databases. *PVLDB* **7**, 277–288 (2013).
13. **Difallah, D. E.**, Demartini, G. & Cudré-Mauroux, P. *Pick-a-crowd: tell me what you like, and i'll tell you what to do* in *22nd International World Wide Web Conference, WWW '13, Rio de Janeiro, Brazil, May 13-17, 2013* (2013), 367–374.
14. **Difallah, D. E.**, Cudré-Mauroux, P. & McKenna, S. A. Scalable Anomaly Detection for Smart City Infrastructure Networks. *IEEE Internet Computing* **17**, 39–47 (2013).
15. Curino, C., **Difallah, D. E.**, Pavlo, A. & Cudré-Mauroux, P. *Benchmarking OLTP/web databases in the cloud: the OLTP-bench framework* in *Proceedings of the Fourth International Workshop on Cloud Data Management, CloudDB 2012, Maui, HI, USA, October 29, 2012* (2012), 17–20.
16. Demartini, G., **Difallah, D. E.** & Cudré-Mauroux, P. *ZenCrowd: leveraging probabilistic reasoning and crowdsourcing techniques for large-scale entity linking* in *Proceedings of the 21st World Wide Web Conference 2012, WWW 2012, Lyon, France, April 16-20, 2012* (2012), 469–478.
17. **Difallah, D. E.**, Demartini, G. & Cudré-Mauroux, P. *Mechanical Cheat: Spamming Schemes and Adversarial Techniques on Crowdsourcing Platforms* in *Proceedings of the First International Workshop on Crowdsourcing Web Search, Lyon, France, April 17, 2012* (2012), 26–30.
18. **Difallah, D. E.**, Benton, R. G., Raghavan, V. V. & Johnsten, T. *FAARM: Frequent Association Action Rules Mining Using FP-Tree* in *Data Mining Workshops (ICDMW), 2011 IEEE 11th International Conference on, Vancouver, BC, Canada, December 11, 2011* (2011), 398–404.